

NetSec_Mid_Sem_Q1

Exported on: 2/23/2026, 11:41:02 PM

User

i will provide you the table if its not visible

```
\begin{table}[h] \centering \renewcommand{\arraystretch}{1.3} \setlength{\tabcolsep}{6pt} \begin{tabular}{|c|*{26}{c|}} \hline \textbf{Index} & 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9 & 10 & 11 & 12 & 13 & 14 & 15 & 16 & 17 & 18 & 19 & 20 & 21 & 22 & 23 & 24 & 25 \\ \hline \textbf{M} & A & B & C & D & E & F & G & H & I & J & K & L & M & N & O & P & Q & R & S & T & U & V & W & X & Y & Z \\ \hline \textbf{E}(\text{textsubscript}{K}(M)) & & & & & & & & & & & & & & & & & & & & & & & & & & & & & & \\ \hline \end{tabular} \caption{Pseudo-random permutation used in block cipher modes} \end{table}
```

answer only for ecb mode : UOMHDJT

Gemini

To decrypt the ciphertext **UOMHDJT** using **Electronic Codebook (ECB) mode**, we simply apply the inverse of the given encryption function to each character individually.

In ECB mode, each block (in this case, each letter) is decrypted independently using the decryption key D_K . Mathematically, for a ciphertext letter C , the plaintext letter P is:

$$P = D_K(C)$$

Step-by-Step Decryption

We look up each letter of the ciphertext in the $E_K(M)$ row of your table and find the corresponding letter in the M row.

Ciphertext (C)	Find in $E_K(M)$	Corresponding M (Plaintext)
U	Index 18	S
O	Index 10	K
M	Index 24	Y
H	Index 19	T
D	Index 0	A
J	Index 11	L
T	Index 4	E

 Export to Sheets

**Final Result**

The decrypted plaintext for the ECB mode ciphertext **UOMHDJT** is: **SKYTALE**

(Note: A "Skytale" or "Scytale" is an ancient Greek tool used for transposition ciphers—a nice cryptographic easter egg in the question!)

Would you like me to proceed with the decryption for the **Cipher Feedback (CFB)** or **Output Feedback (OFB)** modes next?

User

answer for CFB mode: RVPHTUH

Gemini

To decrypt the ciphertext **RVPHTUH** using **Cipher Feedback (CFB)** mode, we need to follow the specific feedback loop where the *previous* ciphertext block is encrypted and then combined with the current ciphertext block.

Rules for this CFB Decryption

Since this is a letter-based cipher, we use the rule provided in your image:

1. The "XOR" operation is replaced by **modulo-26 addition/subtraction**.
2. For decryption, the formula is: $P_i = C_i \ominus E_K(C_{i-1})$ (where \ominus is modulo-26 subtraction).
3. We need an **Initialization Vector (IV)**. Since none is provided in the prompt, it is standard in these textbook problems to assume $IV = 0$ (which corresponds to the letter **A**).

Step-by-Step Calculation

The formula used is: $P_i = (Index(C_i) - Index(E_K(\text{Feedback}))) \pmod{26}$.

Step (<i>i</i>)	Feedback (C_{i-1})	$E_K(\text{Feedback})$	Cipher (C_i)	Calculation (Index)	Plaintext (P_i)
1	A (IV)	D (Index 3)	R (Index 17)	$(17 - 3) = 14$	O
2	R	C (Index 2)	V (Index 21)	$(21 - 2) = 19$	T
3	V	V (Index 21)	P (Index 15)	$(15 - 21) = -6 \equiv 20$	U
4	P	Q (Index 16)	H (Index 7)	$(7 - 16) = -9 \equiv 17$	R
5	H	L (Index 11)	T (Index 19)	$(19 - 11) = 8$	I
6	T	H (Index 7)	U (Index 20)	$(20 - 7) = 13$	N
7	U	B (Index 1)	H (Index 7)	$(7 - 1) = 6$	G

 Export to Sheets

Final Result

The decrypted plaintext for the CFB mode ciphertext **RVPHTUH** is: **OUTRING**

Quick Tip: In CFB decryption, the block cipher (E_K) is still used in its **encryption** direction. We only subtract the result from the ciphertext to recover the plaintext.

Would you like me to move on to the **Output Feedback (OFB)** mode calculation?

User

answer OFB mode: LNMSUU

Gemini

To decrypt the ciphertext **LNMSUU** using **Output Feedback (OFB)** mode, we follow a process similar to CFB, but with one key difference: the feedback loop uses the **output of the encryption function** itself, rather than the previous ciphertext.

Rules for OFB Decryption

- Key Stream Generation:** The key stream Z_i is generated by repeatedly encrypting the previous output: $Z_i = E_K(Z_{i-1})$.
- Plaintext Recovery:** The plaintext is recovered by subtracting the key stream from the ciphertext: $P_i = C_i \ominus Z_i \pmod{26}$.
- Initialization Vector (IV):** Consistent with the previous steps, we assume $IV = 0$ (the letter **A**).

Step-by-Step Calculation

Step (<i>i</i>)	Input to E_K (Z_{i-1})	Key Stream ($Z_i = E_K$)	Cipher (C_i)	Calculation (Index)	Plaintext (P_i)
1	A (IV)	D (Index 3)	L (Index 11)	$(11 - 3) = 8$	I
2	D	X (Index 23)	N (Index 13)	$(13 - 23) = -10 \equiv 16$	Q
3	X	A (Index 0)	M (Index 12)	$(12 - 0) = 12$	M
4	A	D (Index 3)	S (Index 18)	$(18 - 3) = 15$	P
5	D	X (Index 23)	U (Index 20)	$(20 - 23) = -3 \equiv 23$	X
6	X	A (Index 0)	U (Index 20)	$(20 - 0) = 20$	U
7	A	D (Index 3)	Y (Index 24)	$(24 - 3) = 21$	V



Export to Sheets



Final Result

The decrypted plaintext for the OFB mode ciphertext **LNMSUUY** is: **IQMPXUV**

(Note: Unlike the previous two, this appears to be a random string, which is common in cryptographic exercises to ensure you are following the algorithm rather than guessing the word.)

Would you like me to determine the CBC-MAC for the message "TRIPOS" next?

User

determine CBC MAC for TRIPOS (solve correctly)

Gemini

To determine the **CBC-MAC** for the message **TRIPOS**, we use the Cipher Block Chaining (CBC) encryption method, but we only keep the **final** encrypted block as the Message Authentication Code (MAC).

Rules for CBC-MAC

1. **Chaining:** Each plaintext letter is added (modulo-26) to the *previous* encrypted block before being encrypted.
2. **Formula:** $C_i = E_K(P_i \oplus C_{i-1})$.
3. **Initialization:** For a standard MAC, the *IV* is always **0** (which is the letter **A**).
4. **Final Result:** The MAC is the last character C_n .

Step-by-Step Calculation

The message is **T R I P O S**. We process each letter one by one:

Step (i)	Plaintext (P_i)	Prev Cipher (C_{i-1})	Add Mod-26 ($P_i \oplus C_{i-1}$)	Encrypt Result (C_i)
1	T (19)	A (0)	$(19 + 0) = 19 \rightarrow T$	$E_K(T) = H$
2	R (17)	H (7)	$(17 + 7) = 24 \rightarrow Y$	$E_K(Y) = M$
3	I (8)	M (12)	$(8 + 12) = 20 \rightarrow U$	$E_K(U) = B$
4	P (15)	B (1)	$(15 + 1) = 16 \rightarrow Q$	$E_K(Q) = P$
5	O (14)	P (15)	$(14 + 15) = 29 \equiv 3 \rightarrow D$	$E_K(D) = X$
6	S (18)	X (23)	$(18 + 23) = 41 \equiv 15 \rightarrow P$	$E_K(P) = Q$

 Export to Sheets

**Final Result**

The CBC-MAC for the message **TRIPOS** is the final output: **Q** (or Index **16**)

Since you've been working through these cryptographic modes, would you like to explore why **CBC-MAC** is considered insecure for variable-length messages, or should we look at another example?

User

b) Consider another small pseudo-random permutation, this time defined over the set of decimal digits $\{0, 1, 2, \dots, 9\}$, using modulo-10 addition instead of XOR (e.g., $7 \oplus 3 = 0$).

(i) You have intercepted the message 100 with appended CBC-MAC block 4.

The message represents an amount of money to be paid to you and can be of variable length. Use this information to generate a message that represents a much larger number, and provide a valid CBC-MAC digit, without knowing the pseudo-random permutation or key that the recipient will use to verify it. [4 marks]

(ii) What mistake did the designer of the communication system make (leaving aside the tiny block size), and how can

Gemini

Based on the principles of **CBC residues** and **MACs** described in your text, here are the answers to your semester questions.

∅ +2

Part (b)(i): Forging a Larger Message with a Valid CBC-MAC

The goal is to extend the intercepted message $M = (1, 0, 0)$ with MAC $T = 4$ into a longer valid message without knowing the secret key or the encryption function E .

∅ +1

1. Understanding the Intercepted State

In a CBC-MAC, the last block of ciphertext is the MAC. For the message 100 (where each digit is a block), the calculation is:

∅ +2

- $c_1 = E_K(m_1 \oplus IV)$ (Assuming $IV = 0$ for standard MACs).
- $c_2 = E_K(m_2 \oplus c_1)$
- $c_3 = E_K(m_3 \oplus c_2) = 4$ (This is our intercepted MAC).

∅ +1

2. The Forgery Strategy (Length Extension)

We can create a new message M' by appending the intercepted MAC as a "connector" block. Let's construct a 6-digit message: $M' = (1, 0, 0, m_4, m_5, m_6)$. To make the MAC of M' predictable, we want the internal state at the start of the new blocks to effectively "reset" or become known.

∅ +1

- The state after the first three blocks is $c_3 = 4$.
- The next block c_4 is calculated as: $c_4 = E_K(m_4 \oplus c_3)$.

∅ +1

- If we choose m_4 such that it cancels out c_3 , we can reuse the original sequence.
- **Modulo-10 "XOR":** The problem defines \oplus as modulo-10 addition. To "cancel" a 4, we need a number that, when added to 4, results in the original starting block ($m_1 = 1$).
- $m_4 \oplus 4 = 1 \implies m_4 + 4 \equiv 1 \pmod{10} \implies m_4 = 7$.

3. The New Message

- **New Message:** 100700
- **Validation:** * Blocks 1-3 result in $c_3 = 4$.
 - Block 4: $E_K(m_4 \oplus c_3) = E_K(7 + 4 \pmod{10}) = E_K(1) = c_1$.
 - Block 5: $E_K(m_5 \oplus c_4) = E_K(0 \oplus c_1) = c_2$.
 - Block 6: $E_K(m_6 \oplus c_5) = E_K(0 \oplus c_2) = c_3 = 4$.
- **Result:** The message **100700** will have the same valid MAC **4**.

Part (b)(ii): The Designer's Mistake and Fix

The Mistake: Variable Length Support

The designer allowed the message to be of **variable length** without using a length-binding mechanism. Standard CBC-MAC is only secure for messages of a fixed, pre-agreed length. As shown in the attack above, an attacker can append blocks to a known message-MAC pair to create a new valid message because the internal chaining state is exposed by the MAC itself.

🔗 +3

The Fix

To prevent this "length extension" attack, the designer should use one of the following methods:

1. **Prefix the Length:** Include the total message length in the first block before starting the CBC-MAC calculation.
2. **Encrypt the MAC:** Encrypt the final CBC residue one last time with a second, independent secret key.

🔗 +1

3. **XOR the last block:** Use a different key to XOR the final block before the final encryption (a method used in CMAC).

Would you like me to walk through the calculation for a different forged message length, or move on to another question?

Mid Sem NetSec



Show how the DES block cipher can be used to build a 64-bit hash function. Is the result collision resistant?

Building a 64-bit Hash Function Using DES

DES (Data Encryption Standard) is a **64-bit block cipher** with a **64-bit block size** and **56-bit effective key length**. We can construct a hash function from DES using the **Davies-Meyer construction**, which is a common way to build hash functions from block ciphers.

◆ Construction (Davies-Meyer using DES)

Let:

- Message M be divided into 56-bit blocks:

$$M = M_1, M_2, \dots, M_n$$
- H_0 = Initial 64-bit IV (fixed constant)
- $DES_k(x)$ = DES encryption of block x using key k

Since DES key size is 56 bits, each message block is used as the **key**, and the previous hash value is used as the **plaintext**.

For each block:

$$H_i = DES_{M_i}(H_{i-1}) \oplus H_{i-1}$$

Final hash output:

$$H(M) = H_n$$

So the output hash size = **64 bits**

◆ How It Works (Intuition)

1. Start with an initial value H_0
2. Encrypt it using first message block as DES key
3. XOR the output with previous hash value
4. Repeat for all message blocks
5. Final 64-bit value is the hash

This ensures:

- Compression (fixed 64-bit output)
 - Diffusion via DES
 - Dependency on all message blocks
-

◆ Is It Collision Resistant?

 No, it is NOT collision resistant.

Reasons:

1 Small Hash Size (64 bits)

By the **birthday attack**:

- Collision complexity $\approx 2^{n/2}$
- Here $n = 64$

So collision can be found in:

$$2^{32}$$

This is computationally feasible today.

2 DES Itself Is Weak

- Effective key length = 56 bits
- Brute force attack complexity = 2^{56}
- Already practically broken

Example:

DES was broken by the **EFF DES Cracker (1998)** in about 56 hours.

3 Modern Security Standards

Secure hash functions today require:

- At least **128-bit collision resistance**
- Practical minimum output size = **256 bits** (like SHA-256)

Examples:

- MD5 → 128-bit → broken
- SHA-1 → 160-bit → broken
- SHA-256 → secure

◆ Final Conclusion

We can construct a 64-bit hash using DES via the **Davies–Meyer construction**:

$$H_i = DES_{M_i}(H_{i-1}) \oplus H_{i-1}$$

But:

- Output size = **64 bits**
- Collision attack complexity = **2^{32}**
- DES itself is insecure

→ Therefore, the resulting hash function is NOT collision resistant and NOT secure for modern use.

If you want, I can also explain how AES would improve this construction and what output size would be required for real-world security.



An automatic teller machine (ATM) communicates with a central bank computer for PIN verification. A 32-bit CRC code is added to each packet to detect transmission errors and then the link is encrypted using a block cipher in counter mode. Describe an attack that is possible in this setup

🔴 Vulnerability in This ATM Setup

The system does:

1. Append **32-bit CRC** to each packet (for error detection)
2. Encrypt the whole packet using a **block cipher in CTR mode**

At first glance this looks secure — but it is **not**.

🔥 The Core Problem

CTR mode provides:

- Confidentiality
- No Integrity / No Authentication

CRC provides:

- Error detection

- No cryptographic security
- Not keyed
- Linear function

So the system has **no real integrity protection**.

Why This Is Dangerous

CTR mode encryption works as:

$$C = P \oplus K$$

Where:

- P = plaintext
- K = keystream (from counter)
- C = ciphertext

This makes CTR mode:

Fully malleable

An attacker can flip ciphertext bits and control plaintext changes.

Attack: Bit-Flipping Attack (Message Modification)

Step 1: Intercept Packet

Suppose the ATM sends:

```
ini
PIN = 1234
Amount = 1000
```

Packet format:

```
css
[ Data | CRC ]
```

Encrypted using CTR:

$$C = (Data \parallel CRC) \oplus K$$

Attacker intercepts ciphertext.

Step 2: Modify Ciphertext

Attacker flips specific bits in ciphertext.

Because:

$$P = C \oplus K$$

If attacker changes ciphertext:

$$C' = C \oplus \Delta$$

Then plaintext becomes:

$$P' = P \oplus \Delta$$

So attacker can carefully modify:

```
ini
Amount = 1000 → 9000
```

Step 3: Fix the CRC

Here is the critical weakness:

CRC is linear

CRC satisfies:

$$\text{CRC}(x \oplus y) = \text{CRC}(x) \oplus \text{CRC}(y)$$

So attacker can compute:

$$\Delta_{\text{CRC}}$$

and adjust ciphertext so that:

```
kotlin
Modified data + Correct modified CRC
```

Result:

- CRC check passes
- Bank accepts modified transaction

💣 What Can Be Achieved?

An attacker could:

- Change withdrawal amount
- Modify destination account number
- Corrupt transaction fields
- Replay modified packets

This is a **forgery attack**.

🎯 Why This Works

Because:

Mechanism	Protects
CRC	Random noise errors
CTR mode	Confidentiality only
Neither	Message authenticity

There is **no Message Authentication Code (MAC)**.

🔒 What Should Have Been Used?

Instead of:

```
scss
Encrypt( Data || CRC )
```

They should use:

Option 1: Encrypt-then-MAC

```
ini
Ciphertext = Encrypt(Data)
Tag = MAC(Ciphertext)
```

Option 2: Use Authenticated Encryption

- AES-GCM
- AES-CCM

These provide:

- Confidentiality
- Integrity
- Authentication

✓ Final Answer (Exam-Style Summary)

The system is vulnerable to a **bit-flipping attack** because CTR mode encryption is malleable and CRC is linear and not cryptographically secure.

An attacker can modify ciphertext bits to alter the decrypted message and recompute the corresponding CRC difference, causing the modified packet to pass CRC verification.

Thus, the scheme does **not provide integrity or authentication**, and forged transactions are possible.



Consider the problem of mitigation of DDoS attack on a web server. We discussed a solution to this problem using DLP in the class. Propose a solution using Integer Factorization problem.

- a) How do you generate the challenge for the attacker?

- b) How do you generate different challenges for the different attackers?
- c) How do you verify the response given by the attacker?
- d) Explain the total solution along with its merits and demerits

Mitigating DDoS Using the Integer Factorization Problem

(Computational Puzzle-Based Defense)

The idea is to force every client to solve a **computationally expensive problem** before the server allocates heavy resources. Since integer factorization is hard, bots cannot flood the server cheaply.

This is essentially a **Proof-of-Work (PoW)** mechanism based on the difficulty of factoring large composite numbers.

(a) How do you generate the challenge?

The server creates a challenge using the **Integer Factorization Problem**.

Step 1: Generate two primes

Choose large primes:

$$p, q$$

Step 2: Compute composite number

$$N = p \times q$$

Step 3: Send challenge

The challenge to client is:

css

Given N , compute its prime factors p and q

Since integer factorization is hard (basis of RSA security), solving it requires significant computation.

Optional: Include timestamp + nonce to avoid reuse.

(b) How do you generate different challenges for different attackers?

To prevent precomputation and sharing of solutions:

Method 1: Use unique nonce per client

For each connection:

$$N_i = p_i \times q_i$$

where p_i, q_i are freshly generated primes.

Each client receives a different N_i .

Method 2: Bind challenge to client data

Compute:

$$N_i = H(IP \parallel Timestamp) \oplus BaseNumber$$

Then embed this into a composite number construction.

Now challenge depends on:

- Client IP
- Timestamp
- Server secret

So attackers cannot reuse answers.

(c) How do you verify the response?

Client returns:

css

(p, q)

Server verifies by:

$$p \times q = N$$

Verification is **very fast** compared to factoring.

Complexity:

- Factoring → Hard (sub-exponential)
- Multiplication check → $O(n^2)$, very fast

If correct:

- Server allows connection
- Otherwise drop request

(d) Total Solution Explanation

◆ Complete Flow

1. Client requests connection
2. Server detects high load
3. Server sends challenge: factor N
4. Client solves factorization
5. Client sends p, q
6. Server verifies
7. If correct → service granted

◆ Why This Mitigates DDoS

In DDoS:

- Attacker sends millions of requests
- Server resources get exhausted

With factorization puzzle:

- Each request costs attacker heavy CPU
- Server only performs cheap verification

So attack becomes expensive.

This is similar in spirit to **Proof-of-Work systems like** Bitcoin, where computational work prevents spam.

◆ **Merits**

1. Asymmetric cost

- Client work = very high
- Server verification = very low

2. Reduces bot flooding

Bots must spend real computation.

3. No heavy cryptographic state

No complex key management required.

4. Adjustable difficulty

Server can increase bit-length of N during high load.

◆ **Demerits**

1. Legitimate clients also suffer

Mobile devices may struggle.

2. Powerful botnets still possible

If attacker has massive distributed power, attack still works.

3. Prime generation cost on server

Server must generate large primes.

4. Latency increases

Extra round-trip before service.

5. Not suitable for real-time services

◆ **Security Basis**

Security depends on hardness of:

- Integer Factorization Problem

- Same assumption used in RSA

Factoring large (e.g., 1024-bit) composite numbers is computationally hard with classical computers.

✓ Final Summary (Exam-Ready)

A DDoS mitigation scheme can be built using a factorization-based computational puzzle.

The server sends a composite number $N = p \times q$ to the client and asks for its prime factors.

Different challenges are generated per client using fresh primes or client-bound nonces.

Verification is done by checking $p \times q = N$, which is efficient.

The scheme shifts computational burden to the client, limiting large-scale request flooding. However, it increases latency and affects legitimate users.

NetSec_Mid_Sem_Q3

Exported on: 2/24/2026, 12:05:31 AM

User

An automatic teller machine (ATM) communicates with a central bank computer for PIN verification. A 32-bit CRC code is added to each packet to detect transmission errors and then the link is encrypted using a block cipher in counter mode. Describe an attack that is possible in this setup

Gemini

The scenario you've described—using a **32-bit CRC** for integrity and **Counter (CTR) Mode** for encryption—presents a classic vulnerability in cryptographic implementation. While CTR mode provides confidentiality, it is **malleable**, and CRC is **linear**, making the system susceptible to a **Bit-Flipping Attack**.

The Core Vulnerability: Malleability

In CTR mode, encryption is performed by XORing the plaintext (P) with a keystream (K) generated by the block cipher. The ciphertext (C) is:

$$C = P \oplus K$$

Because the keystream is independent of the plaintext, an attacker who knows or can guess parts of the plaintext (like a standard ATM packet header or specific fields) can modify the ciphertext so that it decrypts to a different, predictable plaintext (P').

The Attack Mechanism

The primary goal of the attacker is to modify the packet (e.g., changing a "Transaction Denied" response to "Transaction Approved" or altering the PIN verification result) without the central bank or ATM detecting the change via the 32-bit CRC.

1. Exploiting CTR Mode

An attacker intercepts the ciphertext C . Since XOR is commutative, flipping a bit in the ciphertext results in an identical bit flip in the decrypted plaintext:

$$(P \oplus K) \oplus \text{mask} = (P \oplus \text{mask}) \oplus K$$

If the attacker knows that the 10th bit represents a "Success/Fail" flag, they can simply flip that bit in the ciphertext.

2. Defeating the CRC-32

Usually, a Checksum or CRC is meant to detect these errors. However, **CRC is a linear function**:

$$\text{CRC}(A \oplus B) = \text{CRC}(A) \oplus \text{CRC}(B)$$

This linearity allows the attacker to calculate exactly how the CRC will change when they flip bits in the message.

- The attacker chooses a bitmask (M) to change the message P to P' .
- They calculate $\text{CRC}(M)$.
- They XOR the intercepted ciphertext with M and the intercepted CRC with $\text{CRC}(M)$.
- Upon decryption, the receiver sees the modified message P' , and the modified CRC will match P' perfectly.

The Practical Scenario

Imagine the ATM sends a packet containing: [Account ID | PIN Data | CRC-32]

1. **Intercept:** The attacker sits on the network between the ATM and the Bank.
2. **Modify:** Even if they don't know the PIN, they might know the format of the "Success" packet. They can flip bits to ensure the response sent back from the Bank to the ATM always indicates a successful verification.
3. **Bypass:** Because they can update the CRC-32 mathematically to match their "flipped" bits, the ATM's integrity check passes, and the ATM dispenses cash.

How to Fix It

To prevent this, the system should replace the CRC with a **Message Authentication Code (MAC)** or use an **Authenticated Encryption** mode like **AES-GCM**. Unlike a CRC, a MAC uses a secret key, making it computationally impossible for an attacker to update the integrity tag after modifying the ciphertext.

Would you like me to show you the mathematical proof of why the CRC linearity allows for this specific bit-flipping bypass?

User

The attacker chooses a bitmask (M) to change the message P to P' . They calculate $\text{CRC}(M)$. They XOR the intercepted ciphertext with M and the intercepted CRC with $\text{CRC}(M)$. Upon decryption, the receiver sees the modified message P' , and the modified CRC will match P perfectly.

didn't understand this

Gemini

Think of this like a digital "magic trick" where the attacker changes the contents of a locked box without needing the key.

The reason this attack works is that **CTR mode** (the lock) and **CRC** (the seal) both rely on a mathematical operation called **XOR** (\oplus), which is linear. This allows the attacker to predict exactly how a change to the encrypted data will affect the decrypted result.

Here is the step-by-step breakdown of the logic:

1. The Relationship in CTR Mode

In CTR mode, the ciphertext (C) is just the plaintext (P) mixed with a secret keystream (K):

$$C = P \oplus K$$

If the attacker wants to change the message P to a new message P' , they need a "Difference Mask" (M). This mask represents the bits that are different between the original and the fake message:

$$M = P \oplus P'$$

2. Flipping the Bits

The attacker doesn't know P or K , but they know the **format** (e.g., bit 5 is the "Success" bit). They intercept the ciphertext C and apply the mask M :

$$C_{\text{new}} = C \oplus M$$

When the Bank/ATM decrypts C_{new} , the math looks like this:

$$(P \oplus K) \oplus M = (P \oplus M) \oplus K = P' \oplus K$$

The "mask" passes straight through the encryption and changes the plaintext inside to P' .

3. Fixing the CRC "Seal"

Normally, the CRC would catch this change. However, **CRC is linear**, meaning the CRC of two things added together is the same as the sum of their individual CRCs:

$$\text{CRC}(P \oplus M) = \text{CRC}(P) \oplus \text{CRC}(M)$$

The attacker follows these steps:

1. **Calculate $\text{CRC}(M)$:** Since M is just a string of bits the attacker created, they can calculate its CRC easily.
2. **Modify the Tag:** They take the intercepted CRC from the packet and XOR it with their calculated $\text{CRC}(M)$.
3. **The Result:** When the receiver decrypts the message and checks the CRC, the math "cancels out" perfectly. The modified CRC matches the modified message P' , and the system thinks the data is authentic.

Summary Table: What the Receiver Sees

Component	Original	Attacker's Modification	Receiver's Result
Message	P (Encrypted)	$\oplus M$	P' (Decrypted)
CRC Tag	$CRC(P)$	$\oplus CRC(M)$	$CRC(P')$

 Export to Sheets

Because P' and $CRC(P')$ match, the ATM accepts the fake "Success" command as legitimate.

Would you like me to walk through a simple binary example with a 4-bit message to show how the XOR math cancels out?

User

yes

Gemini

To see how this works, let's use a simplified **4-bit message** instead of a full 32-bit packet. This makes the XOR operations easy to track.

1. The Setup

Assume the ATM is expecting a status bit where **0** means "Denied" and **1** means "Approved."

- **Original Plaintext (P):** `0100` (Let's say the last bit `0` means **Denied**)
- **Secret Keystream (K):** `1011` (Only the Bank and ATM know this)
- **Encryption (CTR Mode):** The ciphertext $C = P \oplus K$
 - `0100` \oplus `1011` = `1111` (This is what the attacker intercepts)

2. The Attacker's Goal

The attacker wants to change the last bit from `0` to `1` so the ATM thinks the transaction is **Approved**.

- **Desired Plaintext (P'):** `0101`
- **Bitmask (M):** The difference between what is there and what they want.
 - $M = P \oplus P' = 0100 \oplus 0101 = 0001$ (A mask to flip the last bit)

3. The CRC Bypass Math

The attacker doesn't know the Keystream (K), but they know how CRC math works. Let's assume a very simple CRC function where $CRC(X)$ is just X itself for this example (in reality, it's a polynomial division, but the **linearity** remains the same).

1. **Intercepted CRC:** $CRC(P) = 0100$
2. **Attacker calculates CRC(M):** $CRC(0001) = 0001$
3. **Attacker modifies the packet:**
 - **New Ciphertext:** $C \oplus M = 1111 \oplus 0001 = 1110$
 - **New CRC Tag:** $CRC(P) \oplus CRC(M) = 0100 \oplus 0001 = 0101$

4. The Receiver (ATM) Decrypts

The ATM receives the modified ciphertext `1110` and the modified CRC `0101`.

- **Step A (Decryption):** The ATM XORs the ciphertext with its secret Keystream (K):
 - `1110` \oplus `1011` = `0101` (The message is now "Approved"!)
- **Step B (Integrity Check):** The ATM calculates the CRC of the decrypted message:
 - $CRC(0101) = 0101$
- **Step C (Verification):** The ATM compares its calculated CRC (`0101`) with the one in the packet (`0101`). **They match!**

The Result

The ATM "verifies" the integrity of the message and finds no errors. It has no way of knowing that the bits were intentionally flipped during transit because the CRC was perfectly updated to match the fake message.

Would you like to see the Python code that simulates this attack to see the XOR logic in action?

User

Answer Q4) (a)

Gemini

To mitigate a DDoS attack using the **Integer Factorization Problem**, we can implement a **Client Puzzle (Proof-of-Work)** mechanism. The goal is to force the client (or attacker) to spend significant computational resources before the server commits its own resources to a request.

Here is how you would generate the challenge for the attacker:

Generating the Challenge

To create a challenge based on integer factorization, the server performs the following steps:

1. **Select Prime Numbers:** The server selects two large, secret prime numbers, p and q .
2. **Compute the Product:** The server computes their product, $n = p \times q$.
3. **Define Difficulty:** The server determines a "difficulty level" by deciding how many bits of the factors to reveal or by choosing the bit-length of the primes.
4. **Issue the Challenge:** The server sends the composite number n to the client. The "challenge" is for the client to find the original prime factors p and q such that $n = p \times q$.

Why this works as a Challenge

- **Asymmetry:** For the server, generating the challenge is extremely "cheap" (just one multiplication: $p \times q$).
- **Computational Burden:** For the attacker, solving the challenge is "expensive." They must use a factorization algorithm (like the General Number Field Sieve) to find the factors.
- **Resource Exhaustion Prevention:** If an attacker tries to flood the server with thousands of requests, they must solve a factorization problem for each one. This slows down the rate at which they can send valid requests, effectively mitigating the DDoS attempt.

Would you like me to move on to part (b) and explain how the server can ensure each attacker receives a unique challenge?

User

part (b)

Gemini

To ensure that an attacker cannot reuse a previously solved puzzle (a "replay attack") or share solutions with other bots in a botnet, the server must generate **unique** challenges for different attackers.

Generating Different Challenges for Different Attackers

The server can achieve this by binding the factorization puzzle to specific attributes of the client's connection. Here are the primary methods:

- **Using Client-Specific Salts:** The server generates primes p and q based on a combination of a master secret and the client's unique identifiers, such as their **IP address** or **Session ID**.
- **Time-Stamping (Nonces):** The server includes a **nonce** (a number used once) or a current timestamp in the challenge generation process. This ensures that even if the same IP makes two requests, the product n will be different each time.
- **Cryptographic Binding:** The server can generate n by hashing the client's metadata (IP + Port + Timestamp) and using that hash as a seed to pick primes from a pre-calculated table or to generate them on the fly.
- **Dynamic Difficulty Scaling:** For different attackers, the server can adjust the "bit-length" of the primes p and q . If a specific IP is flagged as suspicious or is sending requests too frequently, the server can provide that specific attacker with a much larger n , requiring significantly more computational power to factorize.

By incorporating these unique elements, the server ensures that **Puzzle A** solved by **Attacker A** is completely useless for **Attacker B**, forcing every single node in a DDoS swarm to perform its own heavy lifting.

Would you like me to explain part (c), regarding how the server efficiently verifies these responses?

User

part (c) along with merits and demerits of this

Gemini

To complete the solution, we look at the verification phase and evaluate the effectiveness of using the Integer Factorization Problem as a DDoS mitigation tool.

(c) Verifying the Attacker's Response

The beauty of using integer factorization for a client puzzle is that while the **challenge is hard to solve**, the **verification is nearly instantaneous** for the server.

1. **Response Receipt:** The attacker sends back the two factors, p and q .
2. **Simple Multiplication:** The server takes the provided factors and performs a single multiplication: $p \times q$.
3. **Comparison:** The server compares the result of $p \times q$ with the original challenge value n it sent to that specific client.
4. **Granting Access:** * If $(p \times q) == n$, the puzzle is solved. The server accepts the request and processes it.
 - If they do not match, the request is dropped.

Because multiplication is a basic CPU operation (taking nanoseconds), the server can verify thousands of solutions per second without breaking a sweat, while the attacker might have spent seconds or even minutes trying to factorize n .

(d) Merits and Demerits of the Total Solution

Using Integer Factorization as a Proof-of-Work (PoW) mechanism has distinct trade-offs compared to other methods like Discrete Logarithm Problems (DLP) or Hashcash.

Merits

- **Massive Asymmetry:** Factorization is a "One-Way Function." It is computationally easy to multiply but extremely difficult to reverse. This ensures the server stays "relaxed" while the attacker is "exhausted."
- **Granular Control:** The server can easily adjust the difficulty by increasing the bit-length of the primes. A 512-bit n is relatively easy, while a 1024-bit n is exponentially harder.
- **Stateless Potential:** If the server uses a secret key to derive p and q from the client's IP, it doesn't even need to store the factors in a database. it can just re-generate them during verification to save memory.
- **Standardized Math:** Integer factorization is the backbone of RSA security, meaning there are well-optimized libraries available for the server to generate these challenges quickly.

Demerits

- **Hardware Bias:** A legitimate user on a powerful desktop will solve the puzzle much faster than a legitimate user on an old smartphone. This can lead to a poor user experience for people with lower-end hardware.
- **Pre-computation Risk:** If the server doesn't use unique, time-stamped challenges (as discussed in part b), an attacker could build a "Rainbow Table" of factors for common n values.
- **Botnet Scaling:** In a large-scale DDoS attack with millions of compromised devices, even if each bot takes 10 seconds to solve a puzzle, the sheer volume of "solved" requests hitting the server simultaneously can still cause a bottleneck.
- **Energy Inefficiency:** Forcing every legitimate user to run intensive math cycles wastes electricity and battery life, which is why most modern systems prefer "Captcha" or "JS-challenges" over pure mathematical PoW.